

User: alumno  
Passwd: usuario

# Session 1

# Process Management



Departamento de  
Tecnología de los  
Computadores y las  
Comunicaciones  
**Universidad de  
Extremadura**



# Process Management

## Use of Makefiles

```
$ cd ; mkdir ll_1 ; cd ll_1; mkdir calculator ; cd calculator
$ gedit root_of_2.c
```



```
#include <stdio.h>      /* Header of function printf */
#include <stdlib.h>     /* Use of exit() */
#include "calculator.h" /* Header of function square_root */

int main(void)
{
    printf("The square root of 2 is %f.\n", square_root(2.0));
    exit(0);
}
```

root\_of\_2.c



```
$ gedit calculator.c
```



```
#include <math.h>
float add (float x, float y)
{
    return x+y;
}

float square_root (float x)
{
    return sqrt(x);
}
```

Implementation

```
$ gedit calculator.h
```



```
float add      (float x, float y);
float square_root (float x);
```

calculator.h

Interface

calculator.c

How do I  
compile  
this?



**sqrt** is in the mathematical library. That is why we need <math.h>

```
$ gedit Makefile
```



Objective

Precedence

Rule

```
root_of_2: root_of_2.o calculator.o
<TAB>gcc root_of_2.o calculator.o -o root_of_2 -lm

root_of_2.o: root_of_2.c
<TAB>gcc -c root_of_2.c

calculator.o: calculator.c
<TAB>gcc -c calculator.c

clean:
<TAB>rm root_of_2 *.o
```

To link the mathematical library which contains the *sqrt* function

**Makefile**

- Now we have the source files that compose our program. **Next we build the executable.**

- To this end we must create what is known as a **makefile** file.
  - A makefile is a file that provides with a set of **rules** about the steps to generate the final executable
- The left side example contains four rules



A rule is *triggered* if at least one precedence is **more recent** than the objective (see `$ ls -lt`)

```
$ make
gcc -c root_of_2.c
gcc -c calculator.c
gcc -o root_of_2 root_of_2.o calculator.o -lm
```

```
$ ls
calculator.c  calculator.h  calculator.o  makefile  root_of_2
root_of_2.c  root_of_2.o
$ make clean ; ls
calculator.c  calculator.h  makefile  root_of_2.c
$ make
$ ./root_of_2
The square root of 2 is 1.414214.
```

```
$ touch *.c *.h
$ make
```



This touch makes the sources .c and .h more recent than objects .o

## Arguments of a program

We all know that *main* is the first function that will be called inside a program. Its prototype is:


```
int main(int argc, char *argv[]);
```

If we type

```
$ echo Hello world!
```

The **argc** argument passed to the main of echo is **3**

The **argv** passed is as figure shows



```
$ cd .. ; mkdir argv; cd argv
$ gedit arguments.c
$ gcc arguments.c -o arguments
$ ./arguments These are my arguments
```



```
#include <stdio.h>
#include <stdlib.h>

int main(int argc, char *argv[])
{
    while(argc--)
        printf("%s ", *argv++);
    printf("\n");
    exit(EXIT_SUCCESS);
}
```

**arguments.c**

This program prints its arguments

To invoke a system call in your program it is necessary to know:

- Its name, its parameters and its return value.
- Also, the **header** file or files which need to be included in the program.

For instance, if we want to change the permissions of file *bridge.c*, so that every user has read and write access to it, we would invoke the *chmod* system call:

```
#include <sys/types.h> /* This at the beginning of the program */
#include <sys/stat.h>
...
chmod("bridge.c", 0666);
...
```

```
$ man -s 2 chmod
```

```
CHMOD(2)
```

```
Manual del Programador de Linux
```

```
CHMOD(2)
```

```
NOMBRE
```

```
chmod, fchmod - cambia los permisos de un fichero
```

```
SINOPSIS
```

```
#include <sys/types.h>
```

```
#include <sys/stat.h>
```

```
int chmod(const char *path, mode_t mode);
```

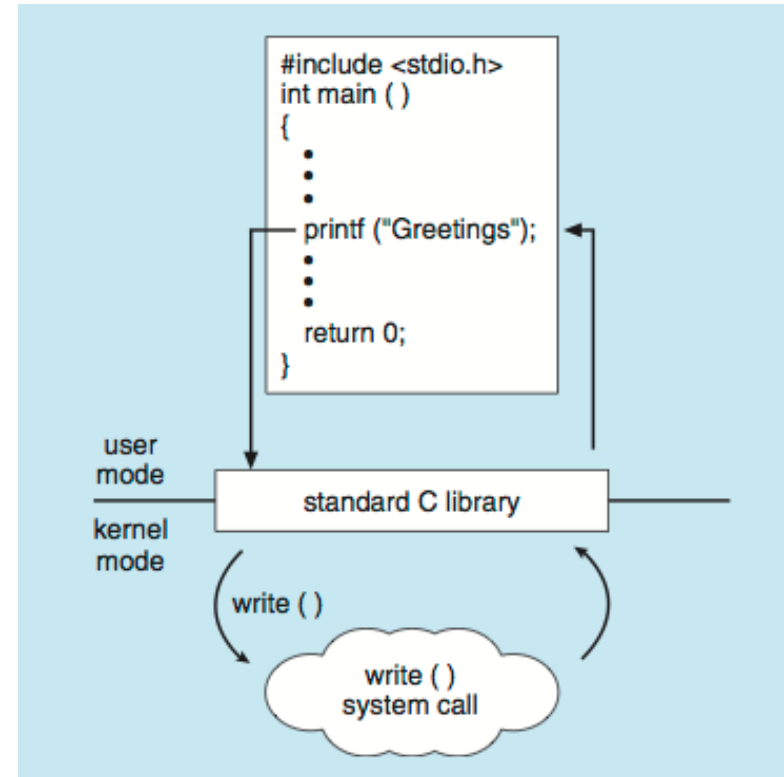
```
int fchmod(int fildes, mode_t mode);
```

```
...
```

```
$ man -s 1 chmod ← This is the chmod command, rather than the system call
```

The manual pages of system calls are in the **section 2** of the manual

- All the functions of the **C library** (printf or malloc, for instance) internally invoke system calls.
- C library function printf, for instance, invokes the write system call, as shown in the figure:



The manual pages of the C library are in the **section 3**

```
$ man -s 3 printf
```

```
PRINTF(3)
```

```
Manual del Programador de Linux
```

```
PRINTF(3)
```

```
NOMBRE
```

```
printf, fprintf, sprintf, vprintf, vfprintf, vsprintf - conversión de salida formateada
```

```
SINOPSIS
```

```
#include <stdio.h>
```

```
int printf(const char *format, ...);
```

```
...
```



## Error conditions

- A system call may fail.  
For instance, the call `chmod` of slide 6 can fail because the file *bridge.c* does not exist, because we have not the suitable permissions to change the permissions of *bridge.c*, etc.
- It is the **responsibility of the programmer** to detect when an error has occurred. We know about the error because the system call returns a special value, usually -1 (minus 1).
- In any case, we should read the manual page to know about this value.
- To detect the error, we would do as follows:

```
#include <sys/types.h>
#include <sys/stat.h>
...
if (0 > chmod("bridge.c",0666)) {
    printf("Can not change the file permissions.\n");
    exit(1);
}
```





## Error conditions

- To know the specific type of error produced we have to read a global variable known as **errno**
- The **manual page** informs on which value errno takes

The manual page of chmod system call says:

```
$ man -s 2 chmod
```

```
...
```

```
ERRORS
```

```
...
```

**EACCES** Search permission is denied on a component of the path prefix.

**EFAULT** Pathname points outside your accessible address space.

**EIO** An I/O error occurred.

**ELOOP** Too many symbolic links were encountered in resolving pathname.

```
...
```

**ENOENT** The file does not exist.

**ENOMEM** Insufficient kernel memory was available.

**ENOTDIR**

```
#include <sys/stat.h>
#include <errno.h>
...
if (0 > chmod("bridge.c", 0666)) {
    if (errno == ENOENT)
        fprintf(stderr, "File bridge.c does not exist.\n");
    else
        fprintf(stderr, "Impossible to change the permissions");
    exit(1);
}
```



## Error conditions

- Considering all and each of the possible value of **errno** is a *costly and tedious task*.
- The useful alternative is the **perror** C library function, which avoids all these tests.
- Function **perror** internally tests the **errno** variable and writes to standard error:
  1. A string that we pass to it as parameter,
  2. colon
  3. and **the system message describing the error**.
- We pass in string (1) the name of the program, the function in the program or any other reference helping the programmer to locate the line which raised the error.
- The former code may appear as follows, by using the function **perror**:

```
#include <sys/stat.h>
```

```
#include <stdio.h>
```

```
...
```

```
if (0 > chmod("bridge.c",0666)) {
```

```
    char my_Message[80];
```

```
    sprintf(my_Message,"%s, chmod, bridge.c", argv[0]);
```

```
    perror(my_Message);
```

```
    exit(1);
```

```
}
```


```
...
```

*Reserve space for the user string*

*Build the user string*

*Pass the user string*

Make sure  
you fully  
understand  
**sprintf**



## Identifiers of users and processes system calls

Calls *getpid*, *getppid* and *getuid* obtain the process identifier (PID) , that of his parent (PPID), and its owner (UID), respectively:

```
#include <sys/types.h>
#include <unistd.h>

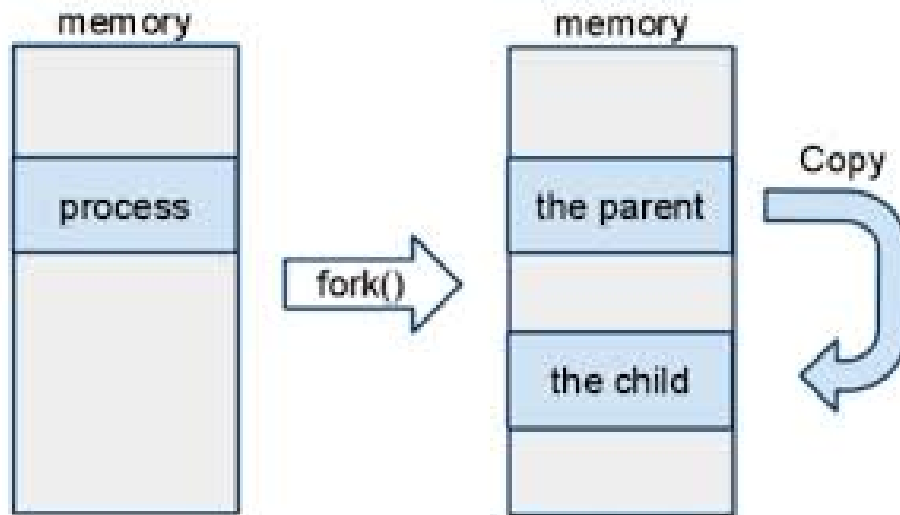
pid_t getpid(void);
pid_t getppid(void);
uid_t getuid(void);
```

**Example** Next program prints the identifier of the invoking process, of its father and its owner.

```
#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>

int main(void)
{
    printf("ID de proceso:          %ld\n", (long)getpid());
    printf("ID de proceso padre:      %ld\n", (long)getppid());
    printf("ID de usuario propietario: %ld\n", (long)getuid());
    return 0;
}
```

- A new process is created by the **fork** system call.
- The new process consists of a copy of the address space of the original process.



- Both processes continue their execution right after the system call `fork()`.
- Since both processes have identical but separate address spaces, those **variables** initialized before the `fork()` call have the same values in both address spaces

### Parent

```
main()
{
    pid = 3456
    pid=fork();
    if (pid == 0)
        ChildProcess();
    else
        ParentProcess();
}

void ChildProcess()
{
    .....
}

void ParentProcess()
{
    .....
}
```

### Child

```
main()
{
    pid = 0
    pid=fork();
    if (pid == 0)
        ChildProcess();
    else
        ParentProcess();
}

void ChildProcess()
{
    .....
}


void ParentProcess()
{
    .....
}
```

Both processes continue execution at the instruction after the fork(), **with one difference**: the return code for the fork() is **zero** for the new (child) process, whereas the **process identifier of the child** is returned to the parent.

## Fork system call

```
#include <sys/types.h>
#include <unistd.h>

pid_t fork(void);
```



**Example** Next program *fork.c* produces a tree like that of the figure:

```
#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>
int main(int argc, char **argv)
{
    int    i;
    pid_t  pid;
    for (i=0; i < 3; i++) {
        if (0 > (pid = fork())) {
            perror("");
            return(1);
        }
        if (pid == 0) {
            fprintf(stdout, "Child: My parent is %ld\n", (long)getppid());
            break;
        }
        fprintf(stdout, "Parent %ld: Created child %ld\n", (long)getpid(), (long)pid);
    }
    return 0;
}
```

```
all: fork
fork: fork.c
<TAB>gcc fork.c -o fork
```

```
clean:
<TAB>rm fork
```

**Makefile**

```
$ cd .. ; mkdir fork; cd fork
$ gedit fork.c
$ gedit Makefile
$ make
$ ./fork
```

**fork.c**

## Exit system call

```
#include <stdlib.h>
void exit (int status);
```

status is between 0 and 255

### Example

```
#include <stdio.h>
#include <stdlib.h>

int main(int argc, char **argv) {

    int dividend = 20;
    int divisor  = atoi(argv[1]);
    int quotient;

    if( divisor == 0) {
        fprintf(stderr, "Division by zero! Exiting...\n");
        exit(1);
    }

    quotient = dividend / divisor;
    fprintf(stderr, "Value of quotient : %d\n", quotient );

    exit(0);
}
```

```
all: fork exit1
fork: fork.c
<TAB>gcc fork.c -o fork
exit1 : exit1.c
<TAB>gcc exit1.c -o exit1

clean:
<TAB>rm fork exit1
```

fork/Makefile

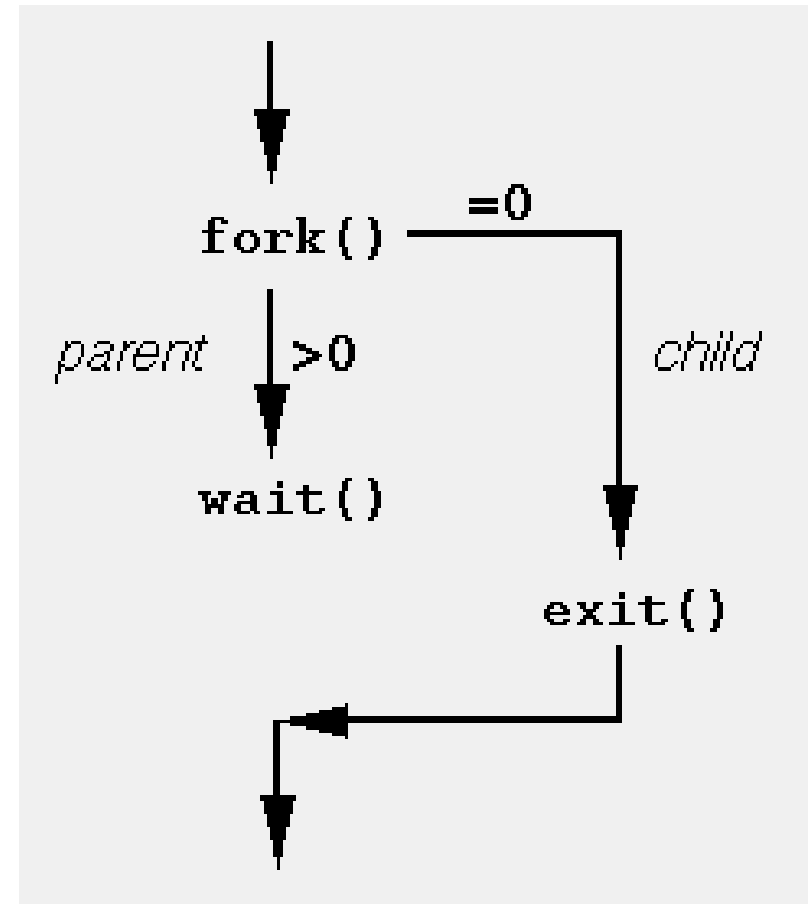
```
$ gedit exit1.c
$ gedit Makefile
$ make
$ ./exit1 2
$ ./exit1 0
```

fork/exit1.c

***Extend!** the fork/Makefile  
with the green code*

A parent can then create more children or, if it has nothing else to do while the child runs, it can issue a **wait** system call to suspend itself until the termination of a child:

```
main()
{
    pid = fork();
    if (pid == 0) {
        /* child */
        ...
        exit(code);
    }
    /* parent */
    wait(&status);
}
```





## Wait system call

```
#include <sys/types.h>
#include <sys/wait.h>
pid_t wait(int *status);
```

status is an output parameter

either the argument of a call to `exit` or the argument of return in `main`

If `status` is different from `NULL`, `wait` fills it with the exit value of the finished child. Once `wait` returns, we can apply macros to status to know precisely the cause of termination of the child:

**WIFEXITED**(status)


If it returns *true* (other than zero), it indicates whether the child ended by calling *exit*.

**WEXITSTATUS**(status)

Returns the *eight* least significant bits of status. Logically, it can only be applied if `WIFEXITED` returns true.

**WIFSIGNALED**(status)

If it returns *true*, the child is terminated because killed by a signal.



**WTERMSIG**(status)

Returns the name of the signal that caused the death of the child process. Logically, this macro can only be applied if `WIFSIGNALED` returned true.

## Wait system call

### Example

```
#include <stdio.h>
#include <unistd.h>
#include <stdlib.h>
#include <sys/wait.h>
int main()
{
    int number, statval;
    printf("%d: I'm the parent !\n", getpid());
    if(fork() == 0) {
        number = 10;
        printf("PID %d: exiting with number %d\n", getpid(), number);
        exit(number);
    }
    else {
        printf("PID %d: waiting for child\n", getpid());
        wait(&statval);
        if(WIFEXITED(statval))
            printf("Child's exit code %d\n", WEXITSTATUS(statval));
        else
            printf("Child did not terminate with exit\n");
    }
    return 0;
}
```

```
all: fork exit1 wait
wait : wait.c
<TAB>gcc wait.c -o wait


clean:
<TAB>rm fork exit1 wait
```

*Keep extending  
the fork/Makefile with  
the green code*

fork/Makefile

```
$ gedit wait.c
$ gedit Makefile
$ make
$ ./wait
```

This should  
return 10

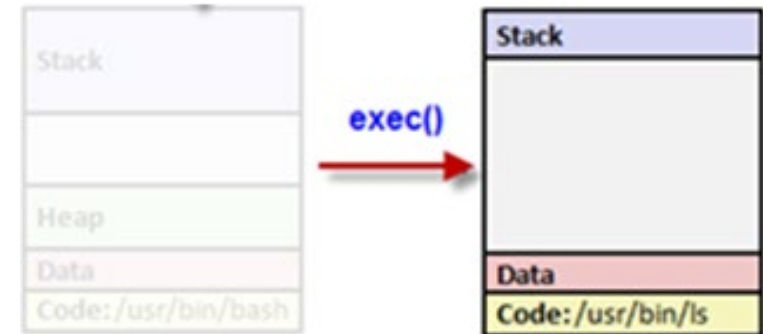


fork/wait.c

## Exec system call

The **exec** family of functions replaces the current process image with a new process image

```
#include <unistd.h>
int execl(char *path,
          char *arg0, ..., char *argn,
          NULL);
```



```
#include <unistd.h>
#include <stdlib.h>
#include <stdio.h>
#include <errno.h>

int main() {
    int status;
    printf ("Mi lista de procesos\n");
    if (0 > execl("/bin/ps", "ps", "ux", NULL)) {
        fprintf(stderr, "Error en exec %d\n", errno);
        exit(1);
    }
    printf ("Fin de mi lista de procesos\n");
    exit(0);
}
```

all: fork exit1 wait **execl**

**execl: execl.c**

<TAB>gcc **execl.c** -o **execl**

clean:

<TAB>rm fork exit1 wait **execl**      **fork/Makefile**


```
$ gedit execl.c
$ gedit Makefile
$ make
$ ./execl
```

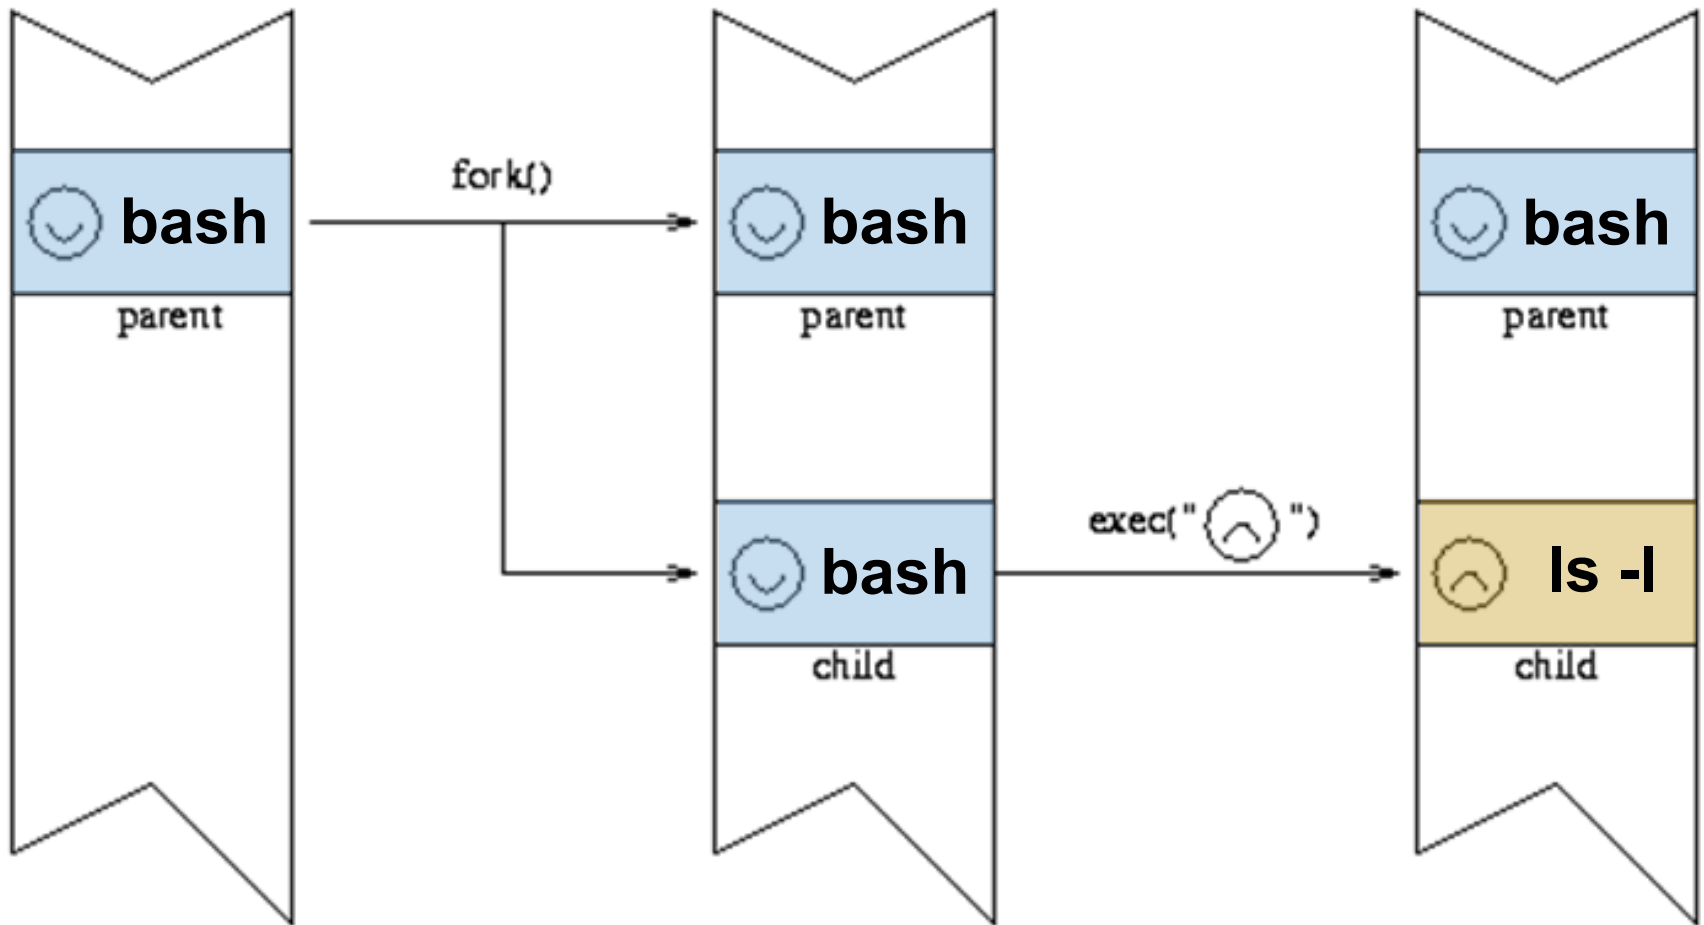


Why does not  
appear on  
the screen?

**execl.c**

After a **fork** system call, the child typically invokes the **exec** system call to replace the process's memory space with a new program.





### Exercise 1

Realice un programa un programa `./exercise1/exhaustFork.c`. Invoca la llamada al sistema `fork` en un bucle hasta que `fork` falla porque se ha terminado la memoria del sistema. Llegado el caso, muestre el número de réplicas que sí han podido crearse. El hijo invoca `sleep(16)` y termina.


### Exercise 2

Realice un programa `./exercise2/launch.c`. Lee de la entrada estándar (mediante la función de la biblioteca C `scanf`) la ruta de un programa y cree un proceso hijo para ejecutar dicho programa. Pruebe a lanzar el programa `arguments.c` de la transparencia 5. Utilice la función `basename` de la biblioteca C para extraer de la ruta del fichero ejecutable el nombre del mismo.

### Home Exercise 3

Realice un programa denominado `./exercise3/four.c`. Creará cuatro procesos, A, B, C y D, de forma que A sea padre de B, B sea padre de C, y C sea padre de D. D debe invocar la llamada al sistema `exec` para mutar en el nuevo proceso `ps`, de modo que la salida en pantalla muestre la relación de parentesco entre los cuatro procesos. Utilice `man ps`.

# References



**The Linux Programming Interface: A Linux and UNIX System Programming Handbook, 1st Edition, [Michael Kerrisk](#)**